

Lecture Slides for

INTRODUCTION TO MACHINE LEARNING

3RD EDITION

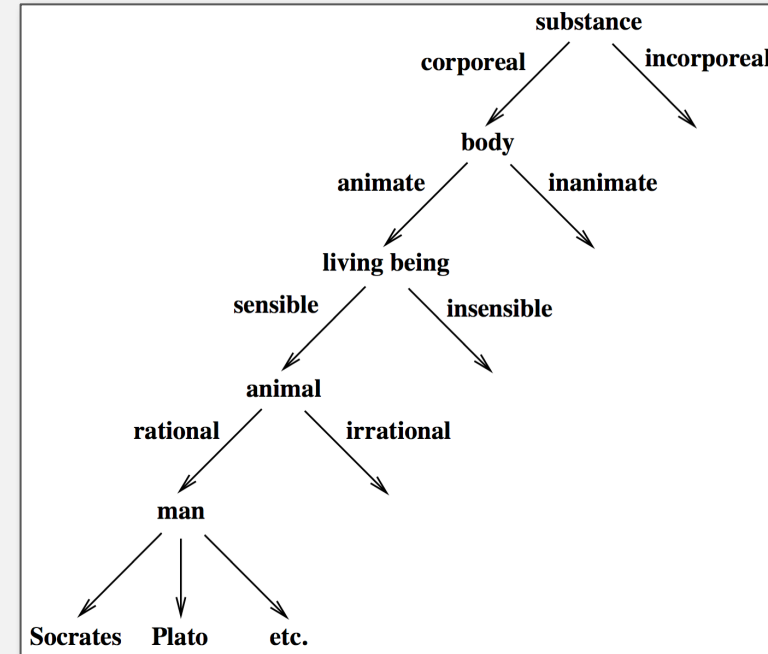

ETHEM ALPAYDIN
© The MIT Press, 2014

alpaydin@boun.edu.tr
<http://www.cmpe.boun.edu.tr/~ethem/i2ml3e>

CHAPTER 9:

DECISION TREES



TREE OF PORPHYRY



- Porphyry of Tyre (c.234 – c.305 CE): Platonic philosopher
- Concept definition by dichotomy (< Plato)

Tree Uses Nodes and Leaves

4



SUITABLE PROBLEMS FOR DECISION TREE LEARNING

- Instances are represented by attribute-value pairs
- Target function has discrete output values
- Disjunctive descriptions may be required
- Training data may contain errors
- Training data may contain missing attribute values

(Mitchell, 1997)

Divide and Conquer

6

- Internal decision nodes
 - ▣ Univariate: Uses a single attribute, x_i
 - Numeric x_i : Binary split : $x_i > w_m$
 - Discrete x_i : n -way split for n possible values
 - ▣ Multivariate: Uses all attributes, \mathbf{x}
- Leaves
 - ▣ Classification: Class labels, or proportions
 - ▣ Regression: Numeric; r average, or local fit
- Learning is **greedy**; find the best split recursively (Breiman et al, 1984; Quinlan, 1986, 1993)

SHANNON INFORMATION (VERY BRIEFLY!)

- Information varies directly with surprise
- Information varies inversely with probability
- Information is additive
- \therefore The information content of a message is proportional to the negative log of its probability

$$I\{s\} = -\lg \Pr\{s\}$$

ENTROPY

- Suppose have source S of symbols from ensemble $\{s_1, s_2, \dots, s_N\}$
- Average information per symbol:

$$\sum_{k=1}^N \Pr\{s_k\} I\{s_k\} = \sum_{k=1}^N \Pr\{s_k\} (-\lg \Pr\{s_k\})$$

- This is the *entropy* of the source:

$$H\{S\} = -\sum_{k=1}^N \Pr\{s_k\} \lg \Pr\{s_k\}$$

MAXIMUM AND MINIMUM ENTROPY

- Maximum entropy is achieved when all signals are equally likely

No ability to guess; maximum surprise


$$H_{\max} = \lg N$$

- Minimum entropy occurs when one symbol is certain and the others are impossible


No uncertainty; no surprise

$$H_{\min} = 0$$


ENTROPY EXAMPLES




$H = 2.0$ bits




$H = 2.0$ bits




$H = 1.9$ bits



$H = 1.0$ bits



$H = 0.3$ bits



$H = 0.0$ bits

Classification Trees (ID3, CART, C4.5)


11

- For node m , N_m instances reach m , N_m^i belong to C_i

$$\hat{P}(C_i | \mathbf{x}, m) \equiv p_m^i = \frac{N_m^i}{N_m}$$

- Node m is **pure** if p_m^i is 0 or 1
- Measure of **impurity** is **entropy**

$$\mathcal{I}_m = -\sum_{i=1}^K p_m^i \log_2 p_m^i$$



SOME IMPURITY MEASURES

- Requirements for 2-class impurity measure:
 $\varphi(p, 1 - p) \geq 0$
 $\varphi(1/2, 1/2) \geq \varphi(p, 1 - p)$ for any $p \in [0, 1]$
 $\varphi(0, 1) = 0 = \varphi(1, 0)$
 $\varphi(p, 1 - p)$ is increasing in p on $[0, 1/2]$ and decreasing in p on $[1/2, 1]$
- Entropy: $\varphi(p, 1 - p) = -p \log_2 p - (1 - p) \log_2 (1 - p)$
- Gini index: $\varphi(p, 1 - p) = 2p(1 - p)$
- Misclassification error: $\varphi(p, 1 - p) = 1 - \max(p, 1 - p)$

Best Split

13

- If node m is pure, generate a leaf and stop, otherwise split and continue recursively
- Impurity after split: N_{mj} of N_m take branch j . N_{mj}^i belong to C_i

$$\hat{p}(C_i | \mathbf{x}, m, j) \equiv p_{mj}^i = \frac{N_{mj}^i}{N_{mj}} \quad \mathcal{I}'_m = - \sum_{j=1}^n \frac{N_{mj}}{N_m} \sum_{i=1}^K p_{mj}^i \log_2 p_{mj}^i$$

- Find the variable and split that minimizes impurity (among all variables — and split positions for numeric variables) \Rightarrow maximum decrease of impurity (greedy)

NUMERIC ATTRIBUTE

- N_m data points reach node m
- There are $N_m - 1$ possible split points
- Best split will always be between values belonging to different classes
- Test halfway points between them to see which leads to highest purity (e.g. minimum entropy)

GenerateTree(\mathcal{X})

If NodeEntropy(\mathcal{X}) $< \theta_I$ /* eq. 9.3

Create leaf labelled by majority class in \mathcal{X}

Return

$i \leftarrow \text{SplitAttribute}(\mathcal{X})$

For each branch of \mathbf{x}_i

Find \mathcal{X}_i falling in branch

GenerateTree(\mathcal{X}_i)

SplitAttribute(\mathcal{X})

MinEnt \leftarrow MAX

For all attributes $i = 1, \dots, d$

If \mathbf{x}_i is discrete with n values

Split \mathcal{X} into $\mathcal{X}_1, \dots, \mathcal{X}_n$ by \mathbf{x}_i

$e \leftarrow \text{SplitEntropy}(\mathcal{X}_1, \dots, \mathcal{X}_n)$ /* eq. 9.8 */

If $e < \text{MinEnt}$ MinEnt $\leftarrow e$; bestf $\leftarrow i$

Else /* \mathbf{x}_i is numeric */

For all possible splits

Split \mathcal{X} into $\mathcal{X}_1, \mathcal{X}_2$ on \mathbf{x}_i

$e \leftarrow \text{SplitEntropy}(\mathcal{X}_1, \mathcal{X}_2)$

If $e < \text{MinEnt}$ MinEnt $\leftarrow e$; bestf $\leftarrow i$

Return bestf

Regression Trees

16

□ Error at node m :

$$b_m(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{x} \in \mathcal{X}_m : \mathbf{x} \text{ reaches node } m \\ 0 & \text{otherwise} \end{cases}$$





$$E_m = \frac{1}{N_m} \sum_t (r^t - g_m)^2 b_m(\mathbf{x}^t) \quad g_m = \frac{\sum_t b_m(\mathbf{x}^t) r^t}{\sum_t b_m(\mathbf{x}^t)}$$

□ After splitting:

$$b_{mj}(\mathbf{x}) = \begin{cases} 1 & \text{if } \mathbf{x} \in \mathcal{X}_{mj} : \mathbf{x} \text{ reaches node } m \text{ and branch } j \\ 0 & \text{otherwise} \end{cases}$$

$$E'_m = \frac{1}{N_m} \sum_j \sum_t (r^t - g_{mj})^2 b_{mj}(\mathbf{x}^t) \quad g_{mj} = \frac{\sum_t b_{mj}(\mathbf{x}^t) r^t}{\sum_t b_{mj}(\mathbf{x}^t)}$$

Model Selection in Trees



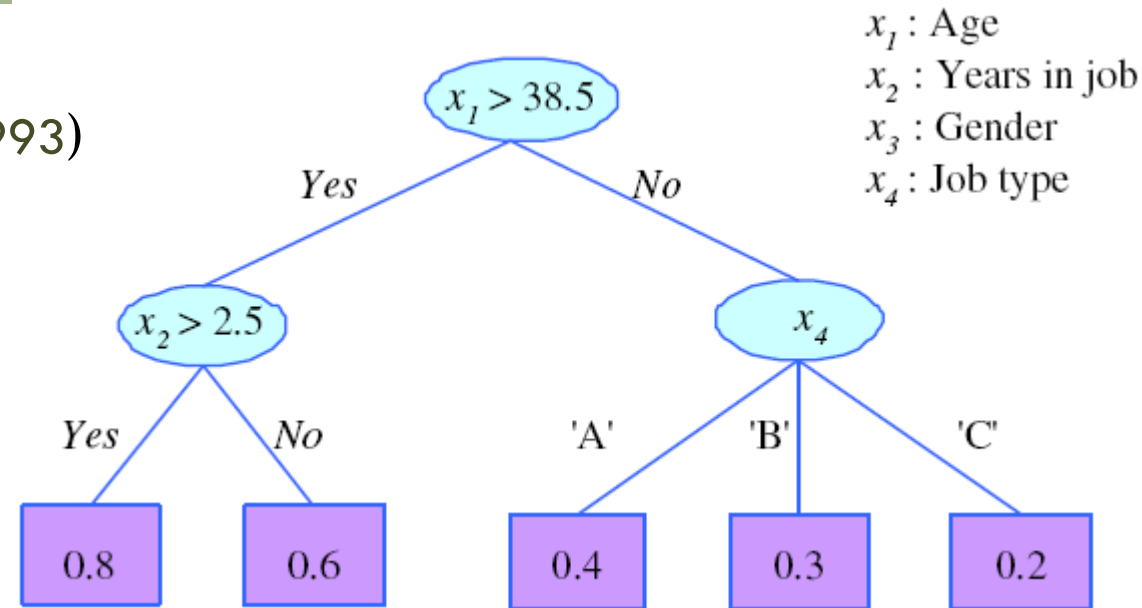
Pruning Trees

18

- Remove subtrees for better generalization (decrease variance)
 - ▣ Prepruning: Early stopping
 - ▣ Postpruning: Grow the whole tree then prune subtrees that overfit on the *pruning set*
- Prepruning is faster, postpruning is more accurate (requires a separate pruning set)

Rule Extraction from Trees

C4.5Rules
(Quinlan, 1993)



- R1: IF (age>38.5) AND (years-in-job>2.5) THEN $y = 0.8$
R2: IF (age>38.5) AND (years-in-job \leq 2.5) THEN $y = 0.6$
R3: IF (age \leq 38.5) AND (job-type='A') THEN $y = 0.4$
R4: IF (age \leq 38.5) AND (job-type='B') THEN $y = 0.3$
R5: IF (age \leq 38.5) AND (job-type='C') THEN $y = 0.2$

RULE POSTPRUNING

1. Infer DT from training set, until training data is fit as well as possible and allowing overfitting to occur
2. Convert learned DT into equivalent rule base (following paths from root to leaves)
3. Prune (generalize) each rule by removing any preconditions that result in improved accuracy on pruning set
4. Sort pruned rules by estimated accuracy, and consider them in this sequence when classifying instances

Learning Rules

21

- Rule induction is similar to tree induction but
 - ▣ tree induction is breadth-first,
 - ▣ rule induction is depth-first; one rule at a time
- Rule set contains rules; rules are conjunctions of terms
- Rule **covers** an example if all terms of the rule evaluate to true for the example
- **Sequential covering**: Generate rules one at a time until all positive examples are covered
 - ▣ Conditions are added one at a time
- IREP (Fürnkranz and Widmer, 1994), Ripper (Cohen, 1995)

```

Ripper(Pos, Neg, k)
  RuleSet  $\leftarrow$  LearnRuleSet(Pos, Neg)
  For  $k$  times
    RuleSet  $\leftarrow$  OptimizeRuleSet(RuleSet, Pos, Neg)
LearnRuleSet(Pos, Neg)
  RuleSet  $\leftarrow \emptyset$ 
  DL  $\leftarrow$  DescLen(RuleSet, Pos, Neg)
  Repeat
    Rule  $\leftarrow$  LearnRule(Pos, Neg)
    Add Rule to RuleSet
    DL'  $\leftarrow$  DescLen(RuleSet, Pos, Neg)
    If DL' > DL + 64
      PruneRuleSet(RuleSet, Pos, Neg)
      Return RuleSet
    If DL' < DL DL  $\leftarrow$  DL'
    Delete instances covered from Pos and Neg
  Until Pos =  $\emptyset$ 
  Return RuleSet

```

```



PruneRuleSet(RuleSet, Pos, Neg)
  For each Rule  $\in$  RuleSet in reverse order
    DL  $\leftarrow$  DescLen(RuleSet, Pos, Neg)
    DL'  $\leftarrow$  DescLen(RuleSet-Rule, Pos, Neg)
    IF DL' < DL Delete Rule from RuleSet
  Return RuleSet

OptimizeRuleSet(RuleSet, Pos, Neg)
  For each Rule  $\in$  RuleSet
    DL0  $\leftarrow$  DescLen(RuleSet, Pos, Neg)
    DL1  $\leftarrow$  DescLen(RuleSet-Rule+
      ReplaceRule(RuleSet, Pos, Neg), Pos, Neg)
    DL2  $\leftarrow$  DescLen(RuleSet-Rule+
      ReviseRule(RuleSet, Rule, Pos, Neg), Pos, Neg)
    If DL1 = min(DL0, DL1, DL2)
      Delete Rule from RuleSet and
      add ReplaceRule(RuleSet, Pos, Neg)
    Else If DL2 = min(DL0, DL1, DL2)
      Delete Rule from RuleSet and
      add ReviseRule(RuleSet, Rule, Pos, Neg)
  Return RuleSet

```

Multivariate Trees

24



DECISION TREES VS. INSTANCE-BASED METHODS

- Each leaf node corresponds to a “bin,” except:
 - bins need not be the same size (as in Parzen windows)
 - bins need not contain equal numbers of training instances (as in kNN)
- Bin divisions are not based only on similarity in input space, but also on supervised output information
 - entropy or mean square error also used
- Thanks to tree structure, leaf (“bin”) is found much faster with smaller number of comparisons
- Decision tree (once constructed) does not store whole training set but only structure of tree, the parameters of the decision nodes, and the output values in leaves
 - implies space complexity is much less than instance-based nonparametric methods (which store all training examples)

(Mitchell, 1997)

FURTHER OBSERVATIONS ON DECISION TREES AND RELATED METHODS

- Considers given features of instances (not some alternative representation)
- Considers the features one at a time
- Sequentially divides the feature space
- Produces explicit rules
- A plausible model of sequential, deliberative, rule-based decision making
- But not of ordinary, real-time cognition and recognition, which is holistic and considers many microfeatures simultaneously

READ ALPAYDIN CH. 10

